

```
pip install tensorflow
```

```
Requirement already satisfied: tensorflow in
/usr/local/lib/python3.11/dist-packages (2.19.0)
Requirement already satisfied: absl-py>=1.0.0 in
/usr/local/lib/python3.11/dist-packages (from tensorflow) (1.4.0)
Requirement already satisfied: astunparse>=1.6.0 in
/usr/local/lib/python3.11/dist-packages (from tensorflow) (1.6.3)
Requirement already satisfied: flatbuffers>=24.3.25 in
/usr/local/lib/python3.11/dist-packages (from tensorflow) (25.2.10)
Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>=0.2.1
in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.6.0)
Requirement already satisfied: google-pasta>=0.1.1 in
/usr/local/lib/python3.11/dist-packages (from tensorflow) (0.2.0)
Requirement already satisfied: libclang>=13.0.0 in
/usr/local/lib/python3.11/dist-packages (from tensorflow) (18.1.1)
Requirement already satisfied: opt-einsum>=2.3.2 in
/usr/local/lib/python3.11/dist-packages (from tensorflow) (3.4.0)
Requirement already satisfied: packaging in
/usr/local/lib/python3.11/dist-packages (from tensorflow) (24.2)
Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!
=4.21.3,!4.21.4,!4.21.5,<6.0.0dev,>=3.20.3 in
/usr/local/lib/python3.11/dist-packages (from tensorflow) (5.29.4)
Requirement already satisfied: requests<3,>=2.21.0 in
/usr/local/lib/python3.11/dist-packages (from tensorflow) (2.32.3)
Requirement already satisfied: setuptools in
/usr/local/lib/python3.11/dist-packages (from tensorflow) (75.2.0)
Requirement already satisfied: six>=1.12.0 in
/usr/local/lib/python3.11/dist-packages (from tensorflow) (1.17.0)
Requirement already satisfied: termcolor>=1.1.0 in
/usr/local/lib/python3.11/dist-packages (from tensorflow) (3.0.1)
Requirement already satisfied: typing-extensions>=3.6.6 in
/usr/local/lib/python3.11/dist-packages (from tensorflow) (4.13.1)
Requirement already satisfied: wrapt>=1.11.0 in
/usr/local/lib/python3.11/dist-packages (from tensorflow) (1.17.2)
Requirement already satisfied: grpcio<2.0,>=1.24.3 in
/usr/local/lib/python3.11/dist-packages (from tensorflow) (1.71.0)
Requirement already satisfied: tensorboard~2.19.0 in
/usr/local/lib/python3.11/dist-packages (from tensorflow) (2.19.0)
Requirement already satisfied: keras>=3.5.0 in
/usr/local/lib/python3.11/dist-packages (from tensorflow) (3.8.0)
Requirement already satisfied: numpy<2.2.0,>=1.26.0 in
/usr/local/lib/python3.11/dist-packages (from tensorflow) (2.0.2)
Requirement already satisfied: h5py>=3.11.0 in
/usr/local/lib/python3.11/dist-packages (from tensorflow) (3.13.0)
Requirement already satisfied: ml-dtypes<1.0.0,>=0.5.1 in
/usr/local/lib/python3.11/dist-packages (from tensorflow) (0.5.1)
Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in
/usr/local/lib/python3.11/dist-packages (from tensorflow) (0.37.1)
Requirement already satisfied: wheel<1.0,>=0.23.0 in
```

```
/usr/local/lib/python3.11/dist-packages (from astunparse>=1.6.0-
>tensorflow) (0.45.1)
Requirement already satisfied: rich in /usr/local/lib/python3.11/dist-
packages (from keras>=3.5.0->tensorflow) (14.0.0)
Requirement already satisfied: namex in
/usr/local/lib/python3.11/dist-packages (from keras>=3.5.0-
>tensorflow) (0.0.8)
Requirement already satisfied: optree in
/usr/local/lib/python3.11/dist-packages (from keras>=3.5.0-
>tensorflow) (0.14.1)
Requirement already satisfied: charset-normalizer<4,>=2 in
/usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0-
>tensorflow) (3.4.1)
Requirement already satisfied: idna<4,>=2.5 in
/usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0-
>tensorflow) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in
/usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0-
>tensorflow) (2.3.0)
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0-
>tensorflow) (2025.1.31)
Requirement already satisfied: markdown>=2.6.8 in
/usr/lib/python3/dist-packages (from tensorboard~=2.19.0->tensorflow)
(3.3.6)
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0
in /usr/local/lib/python3.11/dist-packages (from tensorboard~=2.19.0-
>tensorflow) (0.7.2)
Requirement already satisfied: werkzeug>=1.0.1 in
/usr/local/lib/python3.11/dist-packages (from tensorboard~=2.19.0-
>tensorflow) (3.1.3)
Requirement already satisfied: MarkupSafe>=2.1.1 in
/usr/local/lib/python3.11/dist-packages (from werkzeug>=1.0.1-
>tensorboard~=2.19.0->tensorflow) (3.0.2)
Requirement already satisfied: markdown-it-py>=2.2.0 in
/usr/local/lib/python3.11/dist-packages (from rich->keras>=3.5.0-
>tensorflow) (3.0.0)
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in
/usr/local/lib/python3.11/dist-packages (from rich->keras>=3.5.0-
>tensorflow) (2.19.1)
Requirement already satisfied: mdurl~=0.1 in
/usr/local/lib/python3.11/dist-packages (from markdown-it-py>=2.2.0-
>rich->keras>=3.5.0->tensorflow) (0.1.2)
```

```
import os
import pathlib
import shutil
import random

import tensorflow
```

```

from tensorflow import keras
from tensorflow.keras import layers

import numpy as np
import matplotlib.pyplot as plt

# Detailed shell scripts for setting up datasets
dataset_url =
"https://ai.stanford.edu/~amaas/data/sentiment/aclImdb_v1.tar.gz"
archive_file = "aclImdb_v1.tar.gz"
unsup_folder = "aclImdb/train/unsup"

!curl {dataset_url} -O
!tar -xf {archive_file}
!rm -r {unsup_folder}

```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	
Current			Dload	Upload	Total	Spent	Left
Speed							
100 80.2M	100 80.2M	0 0	18.8M	0	0:00:04	0:00:04	
--:--:--	18.8M						

```

import os

def display_imdb_summary(dataset_path="aclImdb", sample_count=5):
    for data_split in ["train", "test"]:
        print(f"\nSummary of '{data_split}' split:")
        for label in ["pos", "neg"]:
            print(f"    Sentiment: {label}")
            dir_path = os.path.join(dataset_path, data_split, label)
            sample_files = os.listdir(dir_path)[:sample_count]
            for index, sample_name in enumerate(sample_files):
                sample_path = os.path.join(dir_path, sample_name)
                with open(sample_path, "r", encoding="utf-8") as file:
                    sample_lines = file.readlines()
                print(f"    File {index + 1}: {sample_name}")
                print(f"    Lines in file: {len(sample_lines)}")
                print(f"    First 5 lines (or fewer):")
                print("    " + "\n".join(sample_lines[:5]).strip())
            display_imdb_summary()

```

```

Summary of 'train' split:
Sentiment: pos

```

```

File 1: 7880_8.txt
Lines in file: 1

```

First 5 lines (or fewer):

The plot of this enjoyable MGM musical is contrived and only occasionally amusing, dealing with espionage and romance but the focus of the film is properly pointed upon the tuneful interludes showcasing the enormously talented and athletic tap dancing Eleanor Powell, abetted by Tommy Dorsey and his orchestra, featuring Ziggy Elman, Buddy Rich and Frank Sinatra. Red Skelton shares top billing with Powell, and he and sidekick Bert Lahr are given most of the comedic minutes, although Skelton is more effective when he, if it can be believed, performs as Powell's love interest, with Virginia O'Brien actually providing most of the film's humor as the dancer's companion. The technical brilliance of Powell is evidenced during one incredible scene within which Buddy Rich contributes his drumming skills, and which must be viewed several times in order to permit one's breathing to catch up with her precision. Director Edward Buzzell utilizes his large cast well to move the action nicely along despite the rather disjointed script with which he must deal, and permits Powell's cotangent impossibilities to rule the affair, as is appropriate.

File 2: 3815_7.txt

Lines in file: 1

First 5 lines (or fewer):

Corean cinema can be quite surprising for an occidental audience, because of the multiplicity of the tones and genres you can find in the same movie. In a Coreen drama such as this "Secret Sunshine", you'll also find some comical parts, thriller scenes and romantic times. "There's not only tragedy in life, there's also tragic-comedy" says at one point of the movie the character interpreted by Song Kang-ho, summing up the mixture of the picture. But don't get me wrong, this heterogeneity of the genres the movie deals with, adds veracity to the experience this rich movie offers to its spectators. That doesn't mean that it lacks unity : on the contrary, it's rare to see such a dense and profound portrait of a woman in pain.

Shin-ae, who's in quest for a quiet life with her son in the native town of her late husband, really gives, by all the different faces of suffering she's going through, unity to this movie. It's realistic part is erased by the psychological descriptions of all the phases the poor mother is going through. Denial, lost, anger, faith, pert of reality : the movie fallows all the steps the character crosses, and looks like a psychological catalog of all the suffering phases a woman can experience.

The only thing is to accept what may look like a conceptual experience (the woman wears the mask of tragedy, the man represents the comical interludes) and to let the artifices of the movie touch you. I must say that some parts of the movie really did move me (especialy in the beginning), particularly those concerning the inability of Chang Joan to truly help the one he loves, but also that the accumulation of suffering emotionally tired me towards the end. Nevertheless, some cinematographic ideas are really breathtaking and surprising (the scene where a body is discovered in a large shot

is for instance amazing). This kind of scenes makes "Secret Sunshine" the melo equivalent of "The Host" for horror movies or "Memories of murder" for thrillers. These movies are indeed surprising, most original, aesthetically incredible, and manage to give another dimension to the genres they deal with. The only thing that "Secret Sunshine" forgets, as "The host" forgot to be scary, is to make its audience cry : bad point for a melodrama, but good point for a good film.

File 3: 7472_9.txt

Lines in file: 1

First 5 lines (or fewer):

I really liked this movie. I've read a few of the other comments, and although I pity those who did not understand it, I do agree with some of the criticisms. Which, in a strange way, makes me like this movie all the more. I accept that they have got a pretty cast to remake an intelligent movie for the general public, yet it has so many levels and is still great to watch. I also love the movies, such as this one, which provoke so many debates, theories, possible endings and hidden subtext. Congratulations Mr.Crowe, definitely in my Top Ten.

P.S. Saw this when it first came out whilst I was backpacking in Mexico, it was late at night and I had to get back to my hotel and I had a major paranoia trip! Where does the dream end and the real begin?

File 4: 11969_10.txt

Lines in file: 1

First 5 lines (or fewer):

I loved this show growing up and I still watch the first season DVD at age 19 today. What can I say? I grew up in a house much like the one on Full House. I had a dad, two sisters, and a dog. I guess the only difference was that I did not live with my uncle and my dad's best friend. Also, I grew up with my mom in the house. I don't know what I would have done without Full House on television. I think that Stephanie (played by Jodie Sweetin), D.J. (played by Kirk Cameron's sister Candace), and Michelle (Played by Mary-Kate and Ashley Olsen) are my favorite characters. I can relate to each of them because I am the middle child of my family like Steph, I am a younger sister like Michelle, and I am an older sister like D.J. I really like how the show always has moral values because I don't really like any of the O.C.-like shows today. I like the comedy of Full House, too. Uncle Jesse (John Stamos), Joey (Dave Coulier), and Danny (Bob Saget) are hilarious as the girls' uncle, dad's friend, and dad, respectively. The story goes that, after the girls' mom dies, Danny's best friend Joey and his brother-in-law, Jesse move in to help raise the kids. Three men trying to raise three young girls=hilarious. Each character on Full House is full of heart, funny, and genuinely believable. Joey is an aspiring comedian with a kid's heart and soul. Jesse is the cool, motorcycle riding, tough-guy uncle who is softened by his three

nieces, and later, his wife Becky (Laurie Laughlin, from Summerville). Both kids and adults will love this show. Guaranteed.

File 5: 4182_10.txt

Lines in file: 1

First 5 lines (or fewer):

I saw this Australian film about 10 years ago and have never forgotten it. The movie shows the horror of war in a way that Hollywood usually glosses over. The relationship between the soldiers of the two warring countries is highlighted by the differences in culture and the ultimate knowledge that in the end we are all really not different on the inside. If you can find any type of copy of this--buy or rent it. You won't be disappointed, just awed.

Sentiment: neg

File 1: 2308_1.txt

Lines in file: 1

First 5 lines (or fewer):

OK, the box looks promising. Whoopi Goldberg standing next to Danny Glover parodying the famous farmer and his wife painting. Then you pop this baby in the DVD player and all hope is lost in less than five minutes. Supposed to be a comedy. And I must admit I did laugh once about ten minutes before the ending. This movie has the following elements: A battered and abused next door neighbor, a boring legal trial, racism, talk of lynchings, and death and arson. Hilarious, huh? No, please, if you never listen to anyone's reviews, please do here. You cannot even force yourself to watch this crap. CRAP! I said it, CRAP! Whoever put their name on this should indeed sue.

File 2: 10837_1.txt

Lines in file: 1

First 5 lines (or fewer):

I can't tell you how angry I was after seeing this movie. The characters are not the slightest bit interesting, and the plot is nonexistent. So after waiting to see how the lives of these characters affected each other, hoping that the past 2 and a half hours were leading up to some significant finish, what do we get??? A storm of frogs. Now yes, I understand the references to the bible (Exodus) and the underlying theme, but first of all, it was presented with absolutely no resolution, and second of all it would be lost to anyone who has not read the bible (a significant portion of the population) or Charles Fort (a still larger portion). As a somewhat well read person, I thought this movie was a self indulgent poor imitation of a seinfeld episode.

Don't waste your time. It would be better spent reading...

...well anything to be honest

File 3: 890_2.txt

Lines in file: 1

First 5 lines (or fewer):

I saw this regurgitated pile of vignettes tonight at a preview

screening and I was straight up blown away by how bad it was.

First off, the film practically flaunted its gaping blind spots. There are no black or gay New Yorkers in love? Or who, say, know the self-involved white people in love? I know it's not the love Crash of anvil-tastic inclusiveness but you can't pretend to have a cinematic New York with out these fairly prevalent members of society. Plus, you know the people who produced this ish thought Crash deserved that ham-handed Oscar, so where is everyone?

Possibly worse than the bizarre and willful socioeconomic ignorance were the down right offensive chapters (remember when you were in high school and people were openly disgusted with pretty young women in wheelchairs? Me either). This movie ran the gamut of ways to be the worst. Bad acting, bad writing, bad directing -- all spanning every possible genre ever to concern wealthy white people who smoke cigarettes outside fancy restaurants.

But thank god they finally got powerhouses Hayden Christensen and Rachel Bilson back together for that Jumper reunion. And, side note, Uma dodged a bullet; Ethan Hawke looks ravaged. This, of course, is one thing in terms of his looks, but added an incredibly creepy extra vibe of horribleness to his terrifyingly scripted scene opposite poor, lovely Maggie Q.

I had a terrible time choosing my least favorite scene for the end of film questionnaire, but it has to be the Anton Yelchin/Olivia Thirlby bit for the sheer lack of taste, which saddens me because I really like those two actors. I don't consider myself easily offended, but all I could do was scoff and look around with disgust like someone's 50 year old aunt.

A close second place in this incredibly tight contest of terrible things is Shia LaBeouf's tone deaf portrayal of what it means for a former Disney Channel star to act against Julie Christie. I don't mean opposite, I mean against. Against is the only explanation. I realize now that the early sequence with Orlando Bloom is a relative highlight. HIGHLIGHT. Please keep that in mind when your brain begins to leak out your ear soon after the opening credits, which seem to be a nod to the first New York Real World. This film is embarrassing, strangely dated, inarticulate, ineffective, pretentious and, in the end, completely divorced from any real idea of New York at all.

(The extra star is for the Cloris Leachman/ Eli Wallach sequence, as it is actually quite sweet, but it is only one bright spot in what feels like hours of pointless, masturbatory torment.)

File 4: 7974_3.txt

Lines in file: 1

First 5 lines (or fewer):

It's interesting at first. A naive park ranger (Colin Firth) marries a pretty, mysterious woman (Lisa Zane) he's only known for a short time. They seem to be happy, then she disappears without warning. He searches for her and, after a few dead ends, stumbles upon some of her abused childhood and sleazy recent past, which may include criminal activity. And then, it seems the filmmakers didn't know what

to do with the story. The beginning, while not as suspenseful as it sounds, is at least watchable. Then it ceases to be interesting or even make much sense. And the ending is so lame, so dull, and so devoid of any excitement or intelligence, you'll think the screenwriters didn't know what to do with it and got bored trying. What a sorry waste of a good idea!

File 5: 3619_3.txt

Lines in file: 1

First 5 lines (or fewer):

Basically there's a family where a little boy (Jake) thinks there's a zombie in his closet & his parents are fighting all the time.

This movie is slower than a soap opera... and suddenly, Jake decides to become Rambo and kill the zombie.

OK, first of all when you're going to make a film you must Decide if its a thriller or a drama! As a drama the movie is watchable. Parents are divorcing & arguing like in real life. And then we have Jake with his closet which totally ruins all the film! I expected to see a BOOGEYMAN similar movie, and instead i watched a drama with some meaningless thriller spots.

3 out of 10 just for the well playing parents & descent dialogs. As for the shots with Jake: just ignore them.

Summary of 'test' split:

Sentiment: pos

File 1: 9554_10.txt

Lines in file: 1

First 5 lines (or fewer):

Watch out! This is not a gross out comedy like American Pie or you know the comedies of the new millenium. This is a sweet satire of dog shows focused on a bunch of completely different characters. Oh and the commentator is hilarious!

10/10

File 2: 2027_8.txt

Lines in file: 1

First 5 lines (or fewer):

I seen this movie when it came out. I thought what an average movie. I have now realized that this director was ahead of his time. This is a great movie and great soundtrack. I have seen my share of rock films but although this is far from spinal tap (which I did not like)> This film does take us into the life of an 80s rocker wanting to be nothing but. This is nothing more than our inner child wanting to grow up and to be a *ROCK STAR* Yeah I said it. Everyone wants to grow up and be on the spot light(Weather said or not). This movie just puts you in the core of emotions and you can almost feel the excitement of Izzy. I must admit the acting was less par but still the music and story was enough to hold you in to it, till the credits rolled. Worth the watch especially if you are a fan of ye Ole mighty hair bands.

File 3: 9194_10.txt

Lines in file: 1

First 5 lines (or fewer):

Like in "Les amants du Pont-Neuf" two outsiders lives a love story without concessions. The film consists out of a lot of interesting conversation and a lot of sweet moments. The best one comes in a listening booth. They listen to a record together and once in a while they look at each other. They talk, they like each other. She suggests a change in their lives but he is out of hope. The realistic stylestrokes over the realistic (but) emotive dialogs. A really mathematic screenwriter's work for this film. Spanish novel director Jesus Ponce creates one of the most perfect gallery from the latest year of Spanish cinema.

File 4: 7472_9.txt

Lines in file: 1

First 5 lines (or fewer):

I saw this movie yesterday. I must admit - I love it! It's like early Tarantino, but better. Really a must, but... don't show this movie to anyone younger than 18. It's full of blood and sex (rape scene is great :)). Now I'm just waiting for other movies of this director and a DVD release.

File 5: 11969_10.txt

Lines in file: 1

First 5 lines (or fewer):

Scary in places though the effects did leave something to be desired unless you have bad eyesight or are afraid of the dark. However most of the acting was convincing and most of the effects were well done. I thought the creature looked a bit too much like a man in a gorilla suit for my liking. It reminded me of the original pink panther film.

Sentiment: neg

File 1: 4050_3.txt

Lines in file: 1

First 5 lines (or fewer):

Skip all the subjective "this is a great film" reviews and read the IMDB trailer or the back the KINO videobox (which includes both versions of this flick) which I'll paraphrase: "To the tune of sci-fi score by George Antheil, the camera goes on a sleepwalk through B-Movie hell, all photographed by Will Thompson (who did 'Plan 9 from outer space' & 'Maniac')." You don't know whether to laugh AT the film or WITH it. So if you like self-produced B or C-grade noir-wannabe actors and effects with pretensions of surrealism, this could be for you! Otherwise, get a copy of "Screamplay", a modern low-budget expressionist masterpiece.

File 2: 10837_1.txt

Lines in file: 1

First 5 lines (or fewer):

Possibly the worst movie I have ever seen. Pathetic in almost every way.

I threw the DVD straight in the bin - I didn't even think it was fair to give it to the local thrift shop.

The effects are beyond a joke. The dam control room looks like cardboard. The water looks way out of scale with the backgrounds - nothing works.

Then there is the limp plot - about as much depth as a Scooby Doo cartoon.

I couldn't wait for them all to drown.

File 3: 8906_1.txt

Lines in file: 1

First 5 lines (or fewer):

I got hold of this film on DVD with the title Evil Never Sleeps, it gives front cover billing to Carrie Ann Moss, but she plays such a minor character that I didn't really notice her in the film.

I'm afraid that I consider this one of the worst purchases I have ever made. The dialogue was stilted and the delivery wooden, I found the acting to be disconnected from the plot. Graham's performance to me was of someone who's wondering whether she's left the gas on at home.

All in all both my wife and I found this film painful to watch, and it is not a valuable addition to my collection, watch it at your peril, but spending 90 minutes having your fingernails pulled out would probably be a better way to spend your time.

File 4: 8644_4.txt

Lines in file: 1

First 5 lines (or fewer):

William Petersen (that C.S.I guy) has a small uncredited role but it's the best part of the movie. His character comes across smart ass and tough, and it's a fun surprise to see him in this. He has a range that allows him to play just about anything. After his 5 minutes, it goes from looking cool to just nothing much. It leaves you hoping that his character will reappear in the movie but after 20 minutes you give up hope. The movie itself is pretty poor. Worth a watch on TMN or a pick up at the library but not much more. Too much of it reminds you of L.A Confidential except that where that movie starts to get complicated upon itself, this one is so loose, it steers everywhere but where it should. 2 out of 5 stars

File 5: 5276_4.txt

Lines in file: 1

First 5 lines (or fewer):

The German regional-broadcast-station WDR has shown both "The General" and ODC. On Saturday I've seen "The General" and I thought, it wasn't very bad, but not very good too. But yesterday I've seen ODC and I switched it off after about an hour. Although Kevin Spacey was the main actor the movie was totally confusing and seems restless.

"The General" told the story straight and ordered, but ODC just wanted to be cool. There is a reference on the Guy Ritchie Movies "Lock, Stock and Two Smoking Barrels" and "Snatch", but doesn't have the Coolness of these movies.

So, in the end I would rate it 3 of 10!

Get the directories ready for the split validation

batch_size = 32

base_path = pathlib.Path("aclImdb")

validation_path = base_path / "val"

training_path = base_path / "train"

for label in ("neg", "pos"):

os.makedirs(validation_path / label, exist_ok=True)

all_files = os.listdir(training_path / label)

random.Random(1337).shuffle(all_files)

val_count = int(0.2 * len(all_files))

validation_files = all_files[-val_count:]

for file in validation_files:

source_path = training_path / label / file

target_path = validation_path / label / file

if not os.path.exists(target_path):

shutil.move(source_path, target_path)

Loading the datasets from the directories

training_dataset = keras.utils.text_dataset_from_directory(

base_path / "train", batch_size=batch_size

)

validation_dataset = keras.utils.text_dataset_from_directory(

base_path / "val", batch_size=batch_size

)

test_dataset = keras.utils.text_dataset_from_directory(

base_path / "test", batch_size=batch_size

)

Creating a dataset containing only text (without any labels)

text_only_training = training_dataset.map(lambda x, y: x)

Found 25000 files belonging to 2 classes.

Found 5000 files belonging to 2 classes.

Found 25000 files belonging to 2 classes.

Setting the vectorization parameters

sequence_length = 150

vocab_size = 10000

text_vectorizer = layers.TextVectorization(

max_tokens=vocab_size,

```

        output_mode="int",
        output_sequence_length=sequence_length,
    )

    # Using the training text to adjust the vectorizer
    text_vectorizer.adapt(text_only_training)

    # Use the text vectorizer to tokenize datasets
    tokenized_train = training_dataset.map(
        lambda x, y: (text_vectorizer(x), y),
        num_parallel_calls=4).take(100) # Restricting the training
    samples to 100

    tokenized_val = validation_dataset.map(
        lambda x, y: (text_vectorizer(x), y),
        num_parallel_calls=4).take(10000) # Restricting the validation
    samples to 10,000

    tokenized_test = test_dataset.map(
        lambda x, y: (text_vectorizer(x), y),
        num_parallel_calls=4)

    # Establish a simple text classification model with LSTM and embedding
    layers
    input_layer = keras.Input(shape=(None,), dtype="int64")
    embedding_output = layers.Embedding(input_dim=vocab_size,
        output_dim=128)(input_layer)
    lstm_output = layers.Bidirectional(layers.LSTM(32))(embedding_output)
    dropout_output = layers.Dropout(rate=0.3)(lstm_output)
    final_output = layers.Dense(1, activation="sigmoid")(dropout_output)

    sentiment_model = keras.Model(inputs=input_layer,
        outputs=final_output)

    sentiment_model.compile(
        optimizer="rmsprop",
        loss="binary_crossentropy",
        metrics=["accuracy"])

    sentiment_model.summary()
    Model: "functional_3"

```

Layer (type) Param #	Output Shape
input_layer_3 (InputLayer) 0	(None, None)

embedding_2 (Embedding)	(None, None, 128)	
1,280,000		
bidirectional_3 (Bidirectional)	(None, 64)	
41,216		
dropout_3 (Dropout)	(None, 64)	
0		
dense_3 (Dense)	(None, 1)	
65		

Total params: 1,321,281 (5.04 MB)

Trainable params: 1,321,281 (5.04 MB)

Non-trainable params: 0 (0.00 B)

Establish a checkpoint to preserve the model's peak performance
callbacks = [

keras.callbacks.ModelCheckpoint(filepath="embedding_model.keras", save_
best_only=True)
]

Create the model and save the training data for further examination
history_embedded = sentiment_model.fit(
tokenized_train,
validation_data=tokenized_val,
epochs=10,
callbacks=callbacks
)

Epoch 1/10

100/100 ————— 12s 91ms/step - accuracy: 0.5221 - loss:
0.6926 - val_accuracy: 0.5602 - val_loss: 0.7068

Epoch 2/10

100/100 ————— 9s 88ms/step - accuracy: 0.6584 - loss:
0.6201 - val_accuracy: 0.7444 - val_loss: 0.5469

Epoch 3/10

100/100 ————— 9s 87ms/step - accuracy: 0.7822 - loss:
0.4927 - val_accuracy: 0.8024 - val_loss: 0.4571

Epoch 4/10

100/100 ————— 9s 87ms/step - accuracy: 0.8497 - loss:

```
0.3820 - val_accuracy: 0.8214 - val_loss: 0.4142
Epoch 5/10
100/100 _____ 9s 88ms/step - accuracy: 0.8841 - loss:
0.3054 - val_accuracy: 0.8216 - val_loss: 0.4195
Epoch 6/10
100/100 _____ 9s 87ms/step - accuracy: 0.9095 - loss:
0.2499 - val_accuracy: 0.7982 - val_loss: 0.4481
Epoch 7/10
100/100 _____ 9s 87ms/step - accuracy: 0.9353 - loss:
0.1967 - val_accuracy: 0.8286 - val_loss: 0.4474
Epoch 8/10
100/100 _____ 9s 86ms/step - accuracy: 0.9378 - loss:
0.1772 - val_accuracy: 0.8282 - val_loss: 0.4440
Epoch 9/10
100/100 _____ 9s 85ms/step - accuracy: 0.9543 - loss:
0.1520 - val_accuracy: 0.8240 - val_loss: 0.4711
Epoch 10/10
100/100 _____ 8s 84ms/step - accuracy: 0.9685 - loss:
0.1058 - val_accuracy: 0.8220 - val_loss: 0.5582
```

```
# Retrieve metrics from the model's training process
```

```
training_stats = history_embedded.history
```

```
# Preparing the canvas for two plots adjacent to one other
```

```
plt.figure(figsize=(12, 5))
```

```
# Accuracy Plotting for Training and Validation
```

```
plt.subplot(1, 2, 1)
```

```
plt.plot(training_stats['accuracy'], label='Train Accuracy')
```

```
plt.plot(training_stats['val_accuracy'], label='Validation Accuracy')
```

```
plt.title('Accuracy Across Epochs')
```

```
plt.xlabel('Epoch')
```

```
plt.ylabel('Accuracy')
```

```
plt.legend()
```

```
# Plotting Loss for Training and Validation
```

```
plt.subplot(1, 2, 2)
```

```
plt.plot(training_stats['loss'], label='Train Loss')
```

```
plt.plot(training_stats['val_loss'], label='Validation Loss')
```

```
plt.title('Loss Across Epochs')
```

```
plt.xlabel('Epoch')
```


```
plt.ylabel('Loss')
```

```
plt.legend()
```

```
# Presenting everything clean and neatly
```

```
plt.tight_layout()
```

```
plt.show()
```



Retrieve pre-trained GloVe word vectors from the NLP website at Stanford

```
!wget http://nlp.stanford.edu/data/glove.6B.zip
```

Silently extract the downloaded archive

```
!unzip -q glove.6B.zip
```

```
--2025-04-08 15:23:12-- http://nlp.stanford.edu/data/glove.6B.zip
Resolving nlp.stanford.edu (nlp.stanford.edu)... 171.64.67.140
Connecting to nlp.stanford.edu (nlp.stanford.edu)|171.64.67.140|:80...
connected.
```

```
HTTP request sent, awaiting response... 302 Found
```

```
Location: https://nlp.stanford.edu/data/glove.6B.zip [following]
```

```
--2025-04-08 15:23:12-- https://nlp.stanford.edu/data/glove.6B.zip
```

```
Connecting to nlp.stanford.edu (nlp.stanford.edu)|
```

```
171.64.67.140|:443... connected.
```

```
HTTP request sent, awaiting response... 301 Moved Permanently
```

```
Location: https://downloads.cs.stanford.edu/nlp/data/glove.6B.zip
```

```
[following]
```

```
--2025-04-08 15:23:12--
```

```
https://downloads.cs.stanford.edu/nlp/data/glove.6B.zip
```

```
Resolving downloads.cs.stanford.edu (downloads.cs.stanford.edu)...
```

```
171.64.64.22
```

```
Connecting to downloads.cs.stanford.edu (downloads.cs.stanford.edu)|
```

```
171.64.64.22|:443... connected.
```

```
HTTP request sent, awaiting response... 200 OK
```

```
Length: 862182613 (822M) [application/zip]
```

```
Saving to: 'glove.6B.zip.3'
```

```
glove.6B.zip.3      100%[=====>] 822.24M  5.02MB/s   in
2m 39s
```

```
2025-04-08 15:25:52 (5.16 MB/s) - 'glove.6B.zip.3' saved
```

[862182613/862182613]

replace glove.6B.50d.txt? [y]es, [n]o, [A]ll, [N]one, [r]ename:

Use pre-trained GloVe embeddings to set up the matrix

glove_dim = 100

glove_path = "glove.6B.100d.txt"

The GloVe vectors are loaded into a dictionary

glove_index = {}

with open(glove_path, "r", encoding="utf8") as file:

for line in file:

token, vector = line.strip().split(maxsplit=1)

glove_index[token] = np.fromstring(vector, dtype="f", sep=" ")

Create an embedding matrix that matches the language of our model

vocab_list = text_vectorizer.get_vocabulary()

vocab_lookup = {word: idx for idx, word in enumerate(vocab_list)}

embedding_matrix = np.zeros((vocab_size, glove_dim))

for word, idx in vocab_lookup.items():

if idx < vocab_size:

vector = glove_index.get(word)

if vector is not None:

embedding_matrix[idx] = vector

Use preloaded GloVe weights (frozen) to create an embedding layer

glove_embedding = layers.Embedding(

input_dim=vocab_size,

output_dim=glove_dim,

embeddings_initializer=keras.initializers.Constant(embedding_matrix),

trainable=False, *# When training, keep the GloVe weights fixed*

mask_zero=True

)

Use the pretrained embeddings to define the model structure

input_seq = keras.Input(shape=(None,), dtype="int64")

embedded_seq = glove_embedding(input_seq)

lstm_out = layers.Bidirectional(layers.LSTM(32))(embedded_seq)

dropout_out = layers.Dropout(0.3)(lstm_out)

final_output = layers.Dense(1, activation="sigmoid")(dropout_out)

glove_model = keras.Model(inputs=input_seq, outputs=final_output)

Use pre-trained GloVe embeddings to construct and assemble the sentiment model

input_layer = keras.Input(shape=(None,), dtype="int64")

embedded_output = glove_embedding(input_layer)

lstm_output = layers.Bidirectional(layers.LSTM(32))(embedded_output)

dropout_layer = layers.Dropout(0.3)(lstm_output)


```

prediction = layers.Dense(1, activation="sigmoid")(dropout_layer)
pretrained_model = keras.Model(inputs=input_layer, outputs=prediction)

# Using binary classification parameters to compile the model
pretrained_model.compile(
    optimizer="rmsprop",
    loss="binary_crossentropy",
    metrics=["accuracy"]
)

# Printing the model's structure
pretrained_model.summary()

Model: "functional_5"

```

Layer (type) Connected to	Output Shape	Param #
input_layer_5 (InputLayer)	(None, None)	0
embedding_3 (Embedding) input_layer_5[0][0]	(None, None, 100)	1,000,000
not_equal_3 (NotEqual) input_layer_5[0][0]	(None, None)	0
bidirectional_5 embedding_3[1][0], (Bidirectional) not_equal_3[0][0]	(None, 64)	34,048
dropout_5 (Dropout) bidirectional_5[0][0]	(None, 64)	0
dense_5 (Dense) dropout_5[0][0]	(None, 1)	65

Total params: 1,034,113 (3.94 MB)

Trainable params: 34,113 (133.25 KB)

Non-trainable params: 1,000,000 (3.81 MB)

Configure a callback to preserve the optimal GloVe-based model version

```
glove_callbacks = [  
    keras.callbacks.ModelCheckpoint("pretrained_model.keras",  
    save_best_only=True)  
]
```

Use the tokenized data to train the model

```
history_glove = pretrained_model.fit(  
    tokenized_train,  
    validation_data=tokenized_val,  
    epochs=10,  
    callbacks=glove_callbacks  
)
```

Epoch 1/10

100/100 _____ 16s 129ms/step - accuracy: 0.5537 - loss: 0.6958 - val_accuracy: 0.5894 - val_loss: 0.6661

Epoch 2/10

100/100 _____ 12s 124ms/step - accuracy: 0.6172 - loss: 0.6431 - val_accuracy: 0.6640 - val_loss: 0.6177

Epoch 3/10

100/100 _____ 11s 110ms/step - accuracy: 0.6790 - loss: 0.5965 - val_accuracy: 0.6456 - val_loss: 0.6397

Epoch 4/10

100/100 _____ 13s 126ms/step - accuracy: 0.7028 - loss: 0.5684 - val_accuracy: 0.7444 - val_loss: 0.5197

Epoch 5/10

100/100 _____ 12s 124ms/step - accuracy: 0.7392 - loss: 0.5337 - val_accuracy: 0.7556 - val_loss: 0.5136

Epoch 6/10

100/100 _____ 12s 120ms/step - accuracy: 0.7439 - loss: 0.5133 - val_accuracy: 0.7712 - val_loss: 0.4817

Epoch 7/10

100/100 _____ 13s 127ms/step - accuracy: 0.7576 - loss: 0.4942 - val_accuracy: 0.7766 - val_loss: 0.4729

Epoch 8/10

100/100 _____ 11s 110ms/step - accuracy: 0.7706 - loss: 0.4931 - val_accuracy: 0.7578 - val_loss: 0.4943

Epoch 9/10

100/100 _____ 12s 122ms/step - accuracy: 0.7976 - loss: 0.4509 - val_accuracy: 0.7886 - val_loss: 0.4568

Epoch 10/10

```
100/100 _____ 12s 118ms/step - accuracy: 0.8091 - loss: 0.4228 - val_accuracy: 0.7864 - val_loss: 0.4484
```

```
import matplotlib.pyplot as plt
```

```
# Extracting the training history from GloVe-based model  
metrics = history_glove.history
```

```
# Setting up a wide figure with two subplots  
plt.figure(figsize=(12, 5))
```

```
# Plot accuracy: training vs validation
```


```
plt.subplot(1, 2, 1)  
plt.plot(metrics['accuracy'], label='Train Accuracy')  
plt.plot(metrics['val_accuracy'], label='Validation Accuracy')  
plt.title('Accuracy Across Epochs')  
plt.xlabel('Epoch')  
plt.ylabel('Accuracy')  
plt.legend()
```

```
# Plot loss: training vs validation
```

```
plt.subplot(1, 2, 2)  
plt.plot(metrics['loss'], label='Train Loss')  
plt.plot(metrics['val_loss'], label='Validation Loss')  
plt.title('Loss Across Epochs')  
plt.xlabel('Epoch')  
plt.ylabel('Loss')  
plt.legend()
```

```
# Showing both the plots
```

```
plt.tight_layout()  
plt.show()
```



```

import numpy as np
import matplotlib.pyplot as plt
import time

# Defining the sample counts to use for training
sample_counts = [100, 200, 500, 1000]
custom_acc_list = []
glove_acc_list = []

# Setting up the graph
plt.figure(figsize=(12, 6))
plt.title('Accuracy vs Number of Training Samples')
plt.xlabel('Sample Count')
plt.ylabel('Accuracy')
plt.grid(True)

# Loop through various sample sizes
for i, count in enumerate(sample_counts):
    print(f"\n### Training with {count} samples ###\n")

    # Restricting the training set size
    # Replacing 'train_ds' with 'training_dataset' and
    'text_vectorization' with 'text_vectorizer'
    limited_train_data = training_dataset.map(
        lambda x, y: (text_vectorizer(x), y)).take(count)

    # Training model with custom embeddings
    print(f"Training Custom Embedding Model on {count} samples:")
    sentiment_model.fit( # Replacing 'embedding_model' with
'sentiment_model'
        limited_train_data,
        validation_data=tokenized_val, # Replacing 'int_val_ds' with
'tokenized_val'
        epochs=10,
        verbose=1
    )
    # Replacing 'embedding_model' with 'sentiment_model' and
'int_test_ds' with 'tokenized_test'
    acc_custom = sentiment_model.evaluate(tokenized_test, verbose=1)
[1]
    custom_acc_list.append(acc_custom)
    print(f"Custom Embedding Model Accuracy: {acc_custom:.4f}\n")

    # Training the model with pretrained GloVe embeddings
    print(f"Training Pretrained Embedding Model on {count} samples:")
    pretrained_model.fit(
        limited_train_data,
        validation_data=tokenized_val, # Replacing 'int_val_ds' with
'tokenized_val'
        epochs=10,

```

```

        verbose=1
    )
    # Replacing 'int_test_ds' with 'tokenized_test'
    acc_glove = pretrained_model.evaluate(tokenized_test, verbose=1)
[1]
    glove_acc_list.append(acc_glove)
    print(f"Pretrained Embedding Model Accuracy: {acc_glove:.4f}\n")

    # Plotting once all sample sizes are done running
    if i == len(sample_counts) - 1:
        plt.plot(sample_counts, custom_acc_list, marker='o',
label='Custom Embedding', color='blue')
        plt.plot(sample_counts, glove_acc_list, marker='o',
label='Pretrained Embedding', color='orange')

# Final designing of the chart
plt.title('Accuracy vs Number of Training Samples')
plt.xlabel('Sample Count')
plt.ylabel('Accuracy')
plt.xticks(sample_counts)
plt.grid(True)
plt.legend()

# Displaying the chart
plt.tight_layout()
plt.show()

```

Training with 100 samples

Training Custom Embedding Model on 100 samples:

Epoch 1/10

100/100 ————— 9s 85ms/step - accuracy: 0.9605 - loss: 0.1176 - val_accuracy: 0.8344 - val_loss: 0.5587

Epoch 2/10

100/100 ————— 9s 86ms/step - accuracy: 0.9831 - loss: 0.0663 - val_accuracy: 0.7908 - val_loss: 0.5306

Epoch 3/10

100/100 ————— 9s 85ms/step - accuracy: 0.9757 - loss: 0.0799 - val_accuracy: 0.8058 - val_loss: 0.5644

Epoch 4/10

100/100 ————— 9s 85ms/step - accuracy: 0.9797 - loss: 0.0816 - val_accuracy: 0.8138 - val_loss: 0.6496

Epoch 5/10

100/100 ————— 9s 87ms/step - accuracy: 0.9952 - loss: 0.0265 - val_accuracy: 0.8258 - val_loss: 0.6288

Epoch 6/10

100/100 ————— 9s 86ms/step - accuracy: 0.9854 - loss: 0.0489 - val_accuracy: 0.8308 - val_loss: 0.7016

Epoch 7/10

100/100 _____ 9s 87ms/step - accuracy: 0.9977 - loss:
0.0143 - val_accuracy: 0.8126 - val_loss: 0.6292
Epoch 8/10
100/100 _____ 9s 85ms/step - accuracy: 0.9940 - loss:
0.0284 - val_accuracy: 0.8070 - val_loss: 0.6300
Epoch 9/10
100/100 _____ 9s 85ms/step - accuracy: 0.9974 - loss:
0.0124 - val_accuracy: 0.6648 - val_loss: 1.7311
Epoch 10/10
100/100 _____ 9s 86ms/step - accuracy: 0.9778 - loss:
0.0673 - val_accuracy: 0.8170 - val_loss: 0.8740
782/782 _____ 13s 16ms/step - accuracy: 0.7722 - loss:
1.1256
Custom Embedding Model Accuracy: 0.7732

Training Pretrained Embedding Model on 100 samples:

Epoch 1/10
100/100 _____ 11s 110ms/step - accuracy: 0.8149 - loss:
0.4125 - val_accuracy: 0.8026 - val_loss: 0.4364
Epoch 2/10
100/100 _____ 11s 108ms/step - accuracy: 0.8310 - loss:
0.3836 - val_accuracy: 0.7816 - val_loss: 0.4783
Epoch 3/10
100/100 _____ 10s 102ms/step - accuracy: 0.8264 - loss:
0.3850 - val_accuracy: 0.7856 - val_loss: 0.4604
Epoch 4/10
100/100 _____ 11s 109ms/step - accuracy: 0.8385 - loss:
0.3746 - val_accuracy: 0.8086 - val_loss: 0.4367
Epoch 5/10
100/100 _____ 11s 108ms/step - accuracy: 0.8457 - loss:
0.3411 - val_accuracy: 0.7716 - val_loss: 0.4841
Epoch 6/10
100/100 _____ 11s 106ms/step - accuracy: 0.8533 - loss:
0.3330 - val_accuracy: 0.8096 - val_loss: 0.4342
Epoch 7/10
100/100 _____ 11s 109ms/step - accuracy: 0.8762 - loss:
0.3046 - val_accuracy: 0.7984 - val_loss: 0.4629
Epoch 8/10
100/100 _____ 10s 104ms/step - accuracy: 0.8795 - loss:
0.2841 - val_accuracy: 0.8188 - val_loss: 0.4445
Epoch 9/10
100/100 _____ 11s 108ms/step - accuracy: 0.8898 - loss:
0.2796 - val_accuracy: 0.8196 - val_loss: 0.4384
Epoch 10/10
100/100 _____ 11s 108ms/step - accuracy: 0.9020 - loss:
0.2567 - val_accuracy: 0.8200 - val_loss: 0.4386
782/782 _____ 15s 19ms/step - accuracy: 0.7851 - loss:
0.5020
Pretrained Embedding Model Accuracy: 0.7862

Training with 200 samples

Training Custom Embedding Model on 200 samples:

Epoch 1/10

200/200 _____ 15s 72ms/step - accuracy: 0.9581 - loss: 0.1107 - val_accuracy: 0.8378 - val_loss: 0.3845

Epoch 2/10

200/200 _____ 15s 73ms/step - accuracy: 0.9600 - loss: 0.1202 - val_accuracy: 0.8594 - val_loss: 0.3630

Epoch 3/10

200/200 _____ 15s 74ms/step - accuracy: 0.9723 - loss: 0.0923 - val_accuracy: 0.8104 - val_loss: 0.4578

Epoch 4/10

200/200 _____ 14s 72ms/step - accuracy: 0.9774 - loss: 0.0712 - val_accuracy: 0.8454 - val_loss: 0.3930

Epoch 5/10

200/200 _____ 14s 72ms/step - accuracy: 0.9842 - loss: 0.0586 - val_accuracy: 0.8618 - val_loss: 0.4157

Epoch 6/10

200/200 _____ 15s 73ms/step - accuracy: 0.9890 - loss: 0.0368 - val_accuracy: 0.8678 - val_loss: 0.4119

Epoch 7/10

200/200 _____ 14s 72ms/step - accuracy: 0.9918 - loss: 0.0344 - val_accuracy: 0.8442 - val_loss: 0.5465

Epoch 8/10

200/200 _____ 14s 72ms/step - accuracy: 0.9942 - loss: 0.0234 - val_accuracy: 0.8548 - val_loss: 0.4650

Epoch 9/10

200/200 _____ 15s 73ms/step - accuracy: 0.9929 - loss: 0.0272 - val_accuracy: 0.8630 - val_loss: 0.5182

Epoch 10/10

200/200 _____ 15s 73ms/step - accuracy: 0.9959 - loss: 0.0145 - val_accuracy: 0.8064 - val_loss: 0.9331

782/782 _____ 12s 16ms/step - accuracy: 0.7154 - loss: 1.4618

Custom Embedding Model Accuracy: 0.7169

Training Pretrained Embedding Model on 200 samples:

Epoch 1/10

200/200 _____ 18s 92ms/step - accuracy: 0.8846 - loss: 0.2764 - val_accuracy: 0.8146 - val_loss: 0.4202

Epoch 2/10

200/200 _____ 18s 91ms/step - accuracy: 0.8903 - loss: 0.2815 - val_accuracy: 0.8158 - val_loss: 0.4187

Epoch 3/10

200/200 _____ 18s 92ms/step - accuracy: 0.8992 - loss: 0.2687 - val_accuracy: 0.7802 - val_loss: 0.4642

Epoch 4/10

```
200/200 _____ 19s 93ms/step - accuracy: 0.9002 - loss:
0.2532 - val_accuracy: 0.8334 - val_loss: 0.3851
Epoch 5/10
200/200 _____ 19s 94ms/step - accuracy: 0.9118 - loss:
0.2364 - val_accuracy: 0.8340 - val_loss: 0.3855
Epoch 6/10
200/200 _____ 18s 91ms/step - accuracy: 0.9203 - loss:
0.2141 - val_accuracy: 0.8300 - val_loss: 0.3937
Epoch 7/10
200/200 _____ 19s 93ms/step - accuracy: 0.9237 - loss:
0.2073 - val_accuracy: 0.8266 - val_loss: 0.4063
Epoch 8/10
200/200 _____ 19s 93ms/step - accuracy: 0.9322 - loss:
0.1900 - val_accuracy: 0.8266 - val_loss: 0.4018
Epoch 9/10
200/200 _____ 19s 96ms/step - accuracy: 0.9378 - loss:
0.1730 - val_accuracy: 0.8292 - val_loss: 0.4074
Epoch 10/10
200/200 _____ 19s 94ms/step - accuracy: 0.9410 - loss:
0.1605 - val_accuracy: 0.8322 - val_loss: 0.4040
782/782 _____ 15s 19ms/step - accuracy: 0.7891 - loss:
0.4917
Pretrained Embedding Model Accuracy: 0.7894
```

Training with 500 samples

Training Custom Embedding Model on 500 samples:

```
Epoch 1/10
500/500 _____ 33s 65ms/step - accuracy: 0.9501 - loss:
0.1334 - val_accuracy: 0.8994 - val_loss: 0.2711
Epoch 2/10
500/500 _____ 33s 65ms/step - accuracy: 0.9584 - loss:
0.1199 - val_accuracy: 0.9062 - val_loss: 0.2517
Epoch 3/10
500/500 _____ 33s 66ms/step - accuracy: 0.9680 - loss:
0.0919 - val_accuracy: 0.9258 - val_loss: 0.1993
Epoch 4/10
500/500 _____ 33s 65ms/step - accuracy: 0.9810 - loss:
0.0639 - val_accuracy: 0.9344 - val_loss: 0.2010
Epoch 5/10
500/500 _____ 33s 66ms/step - accuracy: 0.9848 - loss:
0.0473 - val_accuracy: 0.9178 - val_loss: 0.2470
Epoch 6/10
500/500 _____ 33s 65ms/step - accuracy: 0.9910 - loss:
0.0321 - val_accuracy: 0.9360 - val_loss: 0.2263
Epoch 7/10
500/500 _____ 33s 65ms/step - accuracy: 0.9943 - loss:
0.0212 - val_accuracy: 0.9300 - val_loss: 0.2450
Epoch 8/10
```



```
500/500 _____ 32s 65ms/step - accuracy: 0.9948 - loss:
0.0187 - val_accuracy: 0.8532 - val_loss: 0.7217
Epoch 9/10
500/500 _____ 33s 65ms/step - accuracy: 0.9956 - loss:
0.0158 - val_accuracy: 0.9394 - val_loss: 0.2903
Epoch 10/10
500/500 _____ 33s 65ms/step - accuracy: 0.9974 - loss:
0.0092 - val_accuracy: 0.9304 - val_loss: 0.2946
782/782 _____ 13s 16ms/step - accuracy: 0.7912 - loss:
0.9677
Custom Embedding Model Accuracy: 0.7924
```

Training Pretrained Embedding Model on 500 samples:

```
Epoch 1/10
500/500 _____ 42s 84ms/step - accuracy: 0.9037 - loss:
0.2363 - val_accuracy: 0.8390 - val_loss: 0.3714
Epoch 2/10
500/500 _____ 41s 83ms/step - accuracy: 0.8955 - loss:
0.2557 - val_accuracy: 0.8506 - val_loss: 0.3417
Epoch 3/10
500/500 _____ 42s 83ms/step - accuracy: 0.8993 - loss:
0.2466 - val_accuracy: 0.8532 - val_loss: 0.3343
Epoch 4/10
500/500 _____ 42s 84ms/step - accuracy: 0.9073 - loss:
0.2327 - val_accuracy: 0.8588 - val_loss: 0.3235
Epoch 5/10
500/500 _____ 43s 85ms/step - accuracy: 0.9161 - loss:
0.2144 - val_accuracy: 0.8620 - val_loss: 0.3180
Epoch 6/10
500/500 _____ 41s 82ms/step - accuracy: 0.9229 - loss:
0.2006 - val_accuracy: 0.8684 - val_loss: 0.3087
Epoch 7/10
500/500 _____ 42s 83ms/step - accuracy: 0.9266 - loss:
0.1914 - val_accuracy: 0.8742 - val_loss: 0.2932
Epoch 8/10
500/500 _____ 43s 85ms/step - accuracy: 0.9315 - loss:
0.1796 - val_accuracy: 0.8704 - val_loss: 0.3075
Epoch 9/10
500/500 _____ 41s 82ms/step - accuracy: 0.9376 - loss:
0.1680 - val_accuracy: 0.8854 - val_loss: 0.2767
Epoch 10/10
500/500 _____ 42s 84ms/step - accuracy: 0.9395 - loss:
0.1546 - val_accuracy: 0.8854 - val_loss: 0.2773
782/782 _____ 15s 19ms/step - accuracy: 0.8221 - loss:
0.4415
Pretrained Embedding Model Accuracy: 0.8230
```

Training with 1000 samples

Training Custom Embedding Model on 1000 samples:

Epoch 1/10

782/782 ————— 50s 64ms/step - accuracy: 0.9841 - loss: 0.0497 - val_accuracy: 0.9410 - val_loss: 0.1570

Epoch 2/10

782/782 ————— 50s 63ms/step - accuracy: 0.9811 - loss: 0.0567 - val_accuracy: 0.9650 - val_loss: 0.1115

Epoch 3/10

782/782 ————— 50s 64ms/step - accuracy: 0.9887 - loss: 0.0365 - val_accuracy: 0.9906 - val_loss: 0.0419

Epoch 4/10

782/782 ————— 49s 63ms/step - accuracy: 0.9923 - loss: 0.0234 - val_accuracy: 0.9822 - val_loss: 0.0518

Epoch 5/10

782/782 ————— 49s 63ms/step - accuracy: 0.9946 - loss: 0.0176 - val_accuracy: 0.9858 - val_loss: 0.0410

Epoch 6/10

782/782 ————— 49s 63ms/step - accuracy: 0.9964 - loss: 0.0121 - val_accuracy: 0.9950 - val_loss: 0.0163

Epoch 7/10

782/782 ————— 49s 63ms/step - accuracy: 0.9979 - loss: 0.0074 - val_accuracy: 0.9946 - val_loss: 0.0194

Epoch 8/10

782/782 ————— 49s 63ms/step - accuracy: 0.9981 - loss: 0.0073 - val_accuracy: 0.9944 - val_loss: 0.0204

Epoch 9/10

782/782 ————— 49s 63ms/step - accuracy: 0.9981 - loss: 0.0074 - val_accuracy: 0.9962 - val_loss: 0.0125

Epoch 10/10

782/782 ————— 49s 62ms/step - accuracy: 0.9989 - loss: 0.0045 - val_accuracy: 0.9976 - val_loss: 0.0076

782/782 ————— 12s 16ms/step - accuracy: 0.8102 - loss: 1.3481

Custom Embedding Model Accuracy: 0.8109

Training Pretrained Embedding Model on 1000 samples:

Epoch 1/10

782/782 ————— 61s 78ms/step - accuracy: 0.9288 - loss: 0.1823 - val_accuracy: 0.8740 - val_loss: 0.2916

Epoch 2/10

782/782 ————— 64s 81ms/step - accuracy: 0.9202 - loss: 0.1988 - val_accuracy: 0.8522 - val_loss: 0.3268

Epoch 3/10

782/782 ————— 64s 82ms/step - accuracy: 0.9292 - loss: 0.1867 - val_accuracy: 0.8984 - val_loss: 0.2487

Epoch 4/10


782/782 ————— 64s 82ms/step - accuracy: 0.9304 - loss: 0.1739 - val_accuracy: 0.8974 - val_loss: 0.2356

Epoch 5/10

```

782/782 _____ 64s 81ms/step - accuracy: 0.9356 - loss:
0.1641 - val_accuracy: 0.8504 - val_loss: 0.3255
Epoch 6/10
782/782 _____ 63s 81ms/step - accuracy: 0.9382 - loss:
0.1585 - val_accuracy: 0.9070 - val_loss: 0.2208
Epoch 7/10
782/782 _____ 63s 81ms/step - accuracy: 0.9447 - loss:
0.1472 - val_accuracy: 0.9138 - val_loss: 0.2127
Epoch 8/10
782/782 _____ 64s 82ms/step - accuracy: 0.9506 - loss:
0.1348 - val_accuracy: 0.9152 - val_loss: 0.2067
Epoch 9/10
782/782 _____ 62s 79ms/step - accuracy: 0.9487 - loss:
0.1334 - val_accuracy: 0.9160 - val_loss: 0.1979
Epoch 10/10
782/782 _____ 62s 79ms/step - accuracy: 0.9525 - loss:
0.1264 - val_accuracy: 0.9330 - val_loss: 0.1720
782/782 _____ 15s 19ms/step - accuracy: 0.8292 - loss:
0.4327
Pretrained Embedding Model Accuracy: 0.8294

```



```

import pandas as pd

# Save the self-trained embedding model's findings
custom_results = {
    "Data Count": sample_counts,
    "Accuracy (Custom Embedding)": custom_acc_list,
}

```

```

# Save the GloVe model's pretrained findings
pretrained_results = {
    "Data Count": sample_counts,
    "Accuracy (Pretrained Embedding)": glove_acc_list,
}

# Combine both sets of results into a single DataFrame
results_table = pd.DataFrame({
    "Data Count": sample_counts,
    "Accuracy (Custom Embedding)": custom_acc_list,
    "Accuracy (Pretrained Embedding)": glove_acc_list
})

# Printing the final comparison table
print("Final Comparison Summary:")
print(results_table)

Final Comparison Summary:
   Data Count  Accuracy (Custom Embedding)  Accuracy (Pretrained
Embedding)
0          100                0.77316
0.78616
1          200                0.71692
0.78936
2          500                0.79244
0.82304
3         1000                0.81092
0.82944

import numpy as np
import matplotlib.pyplot as plt

# Defining the model types
model_labels = ['Custom Embedding', 'Pretrained Embedding']
accuracy_scores = [
    history_embedded.history['val_accuracy'][-1], # Last validation
    history_glove.history['val_accuracy'][-1]     # Last validation
]

# Creating the bar chart
plt.figure(figsize=(8, 6))
plt.bar(model_labels, accuracy_scores, color=['green', 'yellow'],
width=0.5)

# Adding axis labels and chart titles
plt.ylabel('Validation Accuracy')
plt.title('Comparison of Final Validation Accuracy')


```

```

# Displaying values above each bar
for idx, score in enumerate(accuracy_scores):
    plt.text(idx, score + 0.005, f'{score:.4f}', ha='center',
             fontsize=12)

# Rendering the final chart
plt.tight_layout()
plt.show()

```



```

import numpy as np
import matplotlib.pyplot as plt

# Training the sample counts
train_sizes = [100, 200, 500, 1000]
custom_model_scores = [0.75280, 0.77556, 0.80652, 0.81772] # Custom model accuracy
glove_model_scores = [0.78072, 0.80588, 0.81868, 0.82456] # Pretrained model accuracy

# Bar positioning setting up
bar_width = 0.15

```

```

index_positions = np.arange(len(train_sizes)) # Bar centers

# Initializing figure
plt.figure(figsize=(11, 6))

# Plotting the side-by-side bars
plt.bar(index_positions - bar_width / 2, custom_model_scores,
        width=bar_width, label='Custom Embedding', color='skyblue')
plt.bar(index_positions + bar_width / 2, glove_model_scores,
        width=bar_width, label='Pretrained Embedding', color='yellow')

# Adding the line overlays for visual tracking
plt.plot(index_positions - bar_width / 2, custom_model_scores,
        marker='o', color='orange', linestyle='--', label='Custom
Embedding (Line)')
plt.plot(index_positions + bar_width / 2, glove_model_scores,
        marker='o', color='grey', linestyle='--', label='Pretrained
Embedding (Line)')

# labels and title of the charts
plt.xlabel('Sample Size Used for Training', fontsize=12)
plt.ylabel('Accuracy on Test Set', fontsize=12)
plt.title('Test Accuracy by Sample Size and Model Type', fontsize=14)
plt.xticks(index_positions, train_sizes)
plt.legend()

# Annotation of each bar with its accuracy
for i in range(len(train_sizes)):
    plt.text(index_positions[i] - bar_width / 2,
             custom_model_scores[i] + 0.002,
             f'{custom_model_scores[i]:.4f}', ha='center',
             fontsize=10)
    plt.text(index_positions[i] + bar_width / 2, glove_model_scores[i]
             + 0.002,
             f'{glove_model_scores[i]:.4f}', ha='center', fontsize=10)

# Final display of the charts
plt.tight_layout()
plt.show()

```

